

FMN011 – Project 2

I have used Singular value decomposition (SVD) to compress .jpg images. These are stored either as a 2-dimensional matrix representing a grayscale image, or as a 3-dimensional matrix consisting of 3 *component* matrices, each one essentially a grayscale image that represents each one of the colors red, green and blue, and that together make up a color image.

I copied the A matrix from the pdf to a simple text editor for basic reformatting and then into Matlab. I saved it to file as A.mat with the command `save A`. I then wrote a script for creating the grayscale models of the first assignment. This code works from a cleared session, given the file .mat-file containing matrix A.

```
load A;

FOLDER = 'A_matrix';
UNCOMPRESSED_PATH = [FOLDER, '/Uncompressed.jpg'];

FILE_NAME_START = 'A_matrix_grayscale_SVD_compression_rank_';
FIGURE_TITLE_START = 'Matrix A grayscale,';
FIGURE_TITLE_MIDDLE = 'SVD compression rank ';
CAPTION_START = 'Compression ratio ';

RANKS = 1:3;

figure;
set(gcf, 'Position', [0 0 634 638]);

ranks_size = size(RANKS);
subplot_width = (ranks_size(2) + 1) / 2;
subplot(2, subplot_width, 1)

intens(round(A), {FIGURE_TITLE_START, 'No compression'});
imwrite(divide_all_by_largest_element(A), UNCOMPRESSED_PATH);
uncompressed_size = file_size(UNCOMPRESSED_PATH);

%Space allocation
errors = RANKS;
size_ratios = RANKS;

for k = RANKS
    RANK_STRING = int2str(k);
    FULL_TITLE = {FIGURE_TITLE_START, [FIGURE_TITLE_MIDDLE, RANK_STRING]};

    FILE_NAME = create_image_path(FOLDER, [FILE_NAME_START, RANK_STRING]);
    [A_compressed, size_ratios(k), errors(k)] = ...
        compress_and_create_image(A, k, FILE_NAME, uncompressed_size);
```

```

        subplot(2, subplot_width, k+1);
        intens(round(A_compressed), FULL_TITLE);

        FULL_CAPTION = [CAPTION_START, num2str(size_ratios(k), 3)];
        text(0.15, -0.1, FULL_CAPTION, 'Units', 'normalized');
    end

    generate_table(FOLDER, RANKS, size_ratios, errors);

```

Variables in capital letters are used for numerical constants and strings. To achieve logical automatic naming of images, I do lots of string manipulation throughout the code, starting here with the first few lines and further in the `for` loop.

I create a figure which will host the 2×2 subplots, and specify a size, which I extracted using the `get(gcf, 'Position')` function after manually resizing an earlier resultant figure to appropriate size for the report. A similar process is later repeated for each of the assignments.

I plot the uncompressed image in the top-left corner, corresponding to position 1. This is done by `intens`, which converts the matrix to a grayscale plot and adds a title specified as an input parameter.

```

function [] = intens(M, img_title)
    imagesc(M);
    axis off;
    colormap(gray);
    title(img_title);

```

I save the images to file with `imwrite`, which first required me to divide each element with the max value of the entire matrix. This is done by the function `divide_all_by_largest_element`.

```

function M_norm = divide_all_by_largest_element(M)
    max_val = max(M(:));
    M_norm = M / max_val;

```

In the loop I then call `compress_and_create_image` for ranks `1:3`,

```

function [A_compressed, ratio, error] = ...
    compress_and_create_image(A, k, FILE_NAME, uncompressed_size)

A_compressed = svd_compression(A, k);
imwrite(divide_all_by_largest_element(A_compressed), FILE_NAME);

ratio = compression_ratio(uncompressed_size, FILE_NAME);
error = matrix_relative_error(A, A_compressed);

```

where the actual compression is performed by `svd_compression`

```

function y = svd_compression(M, k)
    [U, S, V] = svd(M);

```

```
Y = U(:,1:k)*S(1:k,1:k)*V(:,1:k)';
```

which I based on https://inst.eecs.berkeley.edu/~ee127a/book/login/1_svd_apps_image.html.

The compressed matrix is saved as an image and the compression ratio is computed using the difference in file sizes between the uncompressed and compressed .jpg images. The error is calculated using `matrix_relative_error`.

```
function result = matrix_relative_error(A, A_compressed)
result = norm(A - A_compressed) / norm(A);
```

I then create a subplot and plot the intensity matrix for the compressed matrix. After the loop is finished I put all the error data into a table using `generate_table`, which also saves it to file for easier importing to the report.

```
function [] = generate_table(FOLDER, RANKS, size_ratios, errors)

SIZE_RATIO_STRINGS = arrayfun(@(x) num2str(x, 3), ...
    size_ratios, 'uniformoutput',false);
ERROR_STRINGS = arrayfun(@(x) num2str(x, 2), ...
    errors, 'uniformoutput',false);

data_table = table(RANKS', SIZE_RATIO_STRINGS', ERROR_STRINGS', ...
    'VariableNames',{'SVD_compression_rank' ...
    'Compression_ratio' 'Relative_error'})

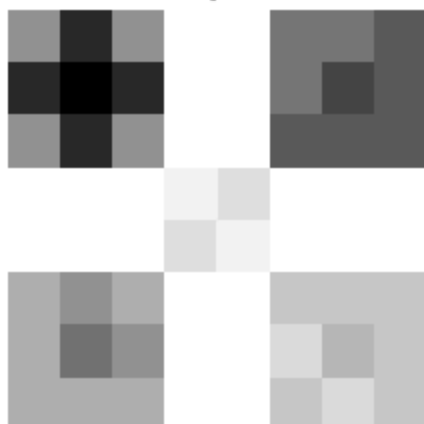
data_table_path = ['data_tables/', FOLDER, '_compression.csv'];
writetable(data_table, data_table_path);
```

The resulting 2×2 subplots are presented below, where it is clear that the compressed image resembles the original more when the rank is increased. Ranks 2 and 3 are *passable* as representations of the uncompressed image, whereas the rank 1 image is compressed to the point where the important contrasts of the image are essentially gone. Maybe if used as a very small thumbnail it could be acceptable.


The compression ratio for each rank is very close to 1, so nothing is really gained from compressing 8×8 pixel images like this, and especially not for ranks above 3 where it is essentially 1. In terms of subjective image quality, the ranks 4? produce excellent representations of the original image, but I couldn't think of a scenario where you would want to sacrifice quality without actually getting a smaller image size.

Dividing every matrix element by the largest element before calling `imagesc` has no effect as far as I can tell, which seems to indicate this function automatically scales its input data.

**Matrix A grayscale,
No compression**



**Matrix A grayscale,
SVD compression rank 1**



Compression ratio 1.05

**Matrix A grayscale,
SVD compression rank 2**



Compression ratio 1.02

**Matrix A grayscale,
SVD compression rank 3**



Compression ratio 1.01

For the the second assignment I started by downloading `nena.jpg` from the web page, and then wrote scripts for the grayscale and color image compressions.

```
COLOR_UNCOMPRESSED_PATH = 'nena.jpg';

nena = imread(COLOR_UNCOMPRESSED_PATH);
file = dir(COLOR_UNCOMPRESSED_PATH);
nena_size = file.bytes;

[grayscale_ranks, grayscale_size_ratios, grayscale_errors] = ...
    assignment_2_grayscale(nena);
[color_ranks, color_size_ratios, color_errors] = ...
    assignment_2_color(nena, nena_size);
```

I calculate the file size for use in computing the compression ratios later on, and call functions `assignment_2_grayscale` and `assignment_2_color`.

```
function [RANKS, size_ratios, errors] = assignment_2_grayscale(nena)

FOLDER = 'nena_grayscale';
UNCOMPRESSED_IMAGE_PATH = [FOLDER, ...
    '/Nena_grayscale_uncompressed.jpg'];
FILE_NAME_START = 'Nena_grayscale_SVD_compressed_rank_';

FIGURE_TITLE = 'Nena grayscale, SVD compression';
SUBPLOT_TITLE_START = 'Rank ';
SUBPLOT_CAPTION_START = 'Compression';
SUBPLOT_CAPTION_MIDDLE = 'ratio ';

RANKS = 2 .^ (0:7);

nena_grayscale = double(rgb2gray(nena));
imwrite(divide_all_by_largest_element(nena_grayscale),
    UNCOMPRESSED_IMAGE_PATH);
uncompressed_size = file_size(UNCOMPRESSED_IMAGE_PATH);

%Space allocation
rank_size = size(RANKS);
rank_indices = 1:rank_size(2);
size_ratios = RANKS;
errors = RANKS;

figure;
set(gcf, 'Position', [0 0 926 597]);

for k = rank_indices
    RANK_STRING = int2str(RANKS(k));
    FILE_NAME = create_image_path(FOLDER, strcat( ...
        FILE_NAME_START, RANK_STRING));
    FULL_TITLE = [SUBPLOT_TITLE_START, RANK_STRING];

    [compressed_image, size_ratios(k), errors(k)] = ...
        compress_and_create_image( nena_grayscale, ...
            RANKS(k), FILE_NAME, uncompressed_size);
```

```

subplot(2, rank_size(2)/2, k);
intens(compressed_image, FULL_TITLE);

FULL_CAPTION = {SUBPLOT_CAPTION_START, ...
    [SUBPLOT_CAPTION_MIDDLE, num2str(size_ratios(k), 3)]};
text(0.2, -0.12, FULL_CAPTION, 'Units', 'normalized');
end

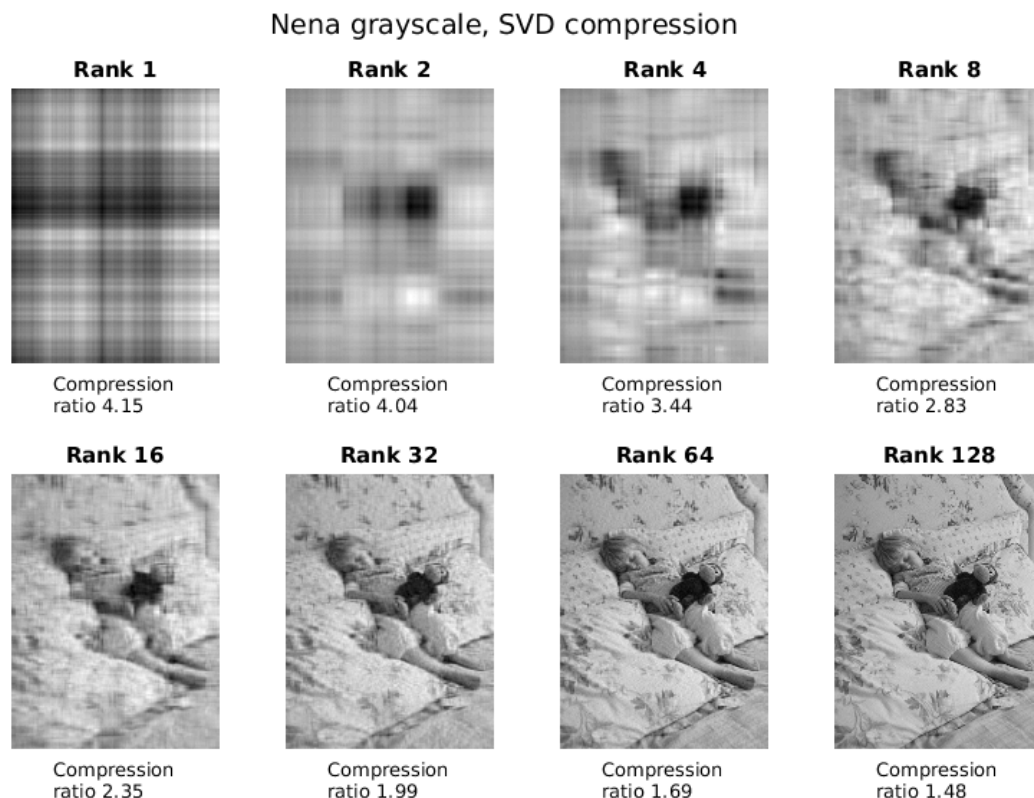
suptitle(FIGURE_TITLE)

generate_table(FOLDER, RANKS, size_ratios, errors);

```

I found that exponential values for RANKS gives a good span and clearly brings out the difference between compression ranks. For each rank specified, `compress_and_create_image` is once again called to compress the matrix and save it as an image. Most of the other code resembles that of the first assignment.

For compression ranks below 2^4 , images are mainly noise and basically useless for any application. At around 2^5 , reasonable quality is achieved for the first time at a compression ratio close to 2. For a maximized image on my regular-sized laptop screen, great quality is given at 2^8 . For even higher ranks, the compressed images look very close to the original, but at the cost of any real decrease in image size.



The color image manipulation is performed by `assignment_2_color`.

```
function [RANKS, size_ratios, errors] = ...
    assignment_2_color(nena, nena_size)

FOLDER = 'nena_color';

FIGURE_TITLE = 'Nena color, SVD compression';
SUBPLOT_TITLE_START = 'Rank ';
SUBPLOT_CAPTION_START = 'Compression';
SUBPLOT_CAPTION_MIDDLE = 'ratio ';

RANKS = 2 .^ (0:7);

nena_double = double(nena);

%Space allocation
ranks_size = size(RANKS);
rank_indices = 1:ranks_size(2);
size_ratios = rank_indices;
errors = rank_indices;

loops = ranks_size(2);
F(loops) = struct('cdata', [], 'colormap', []);

figure;
set(gcf, 'Position', [0 0 926 597]);

for k = rank_indices
    [compressed_tensor, compressed_tensor_divided_by_largest, ...
     size_ratios(k), errors(k)] = compress_3_components(FOLDER, ...
     nena_double, nena_size, RANKS(k));

    movie_image = imresize(uint8(round(compressed_tensor)), 0.3);
    F(k) = im2frame(movie_image);

    subplot(2, ranks_size(2) / 2, k);
    imshow(imresize(compressed_tensor_divided_by_largest, 0.1));

    axis image;
    title([SUBPLOT_TITLE_START, int2str(RANKS(k))]);

    FULL_CAPTION = {SUBPLOT_CAPTION_START, ...
        [SUBPLOT_CAPTION_MIDDLE, num2str(size_ratios(k), 3)]};
    text(0.2, -0.12, FULL_CAPTION, 'Units', 'normalized');
end

suptitle(FIGURE_TITLE);
generate_table(FOLDER, RANKS, size_ratios, errors);

assignment_2_video(F, RANKS);
```

The compression of a color image is more complicated than for grayscale, and is handled by the function `compress_3_components` which uses `svd_compression` on each component

separately and then recombines the compressed components. After this, the image is saved to file, and compression ratios and errors are returned. The errors for the color images are calculated as the mean of the relative errors of each individual component matrix.

```
function [compressed_image, compressed_image_divided_by_largest, ...
    size_ratio, error_median] = compress_3_components(FOLDER, ...
    uncompressed_image, uncompressed_size, rank)

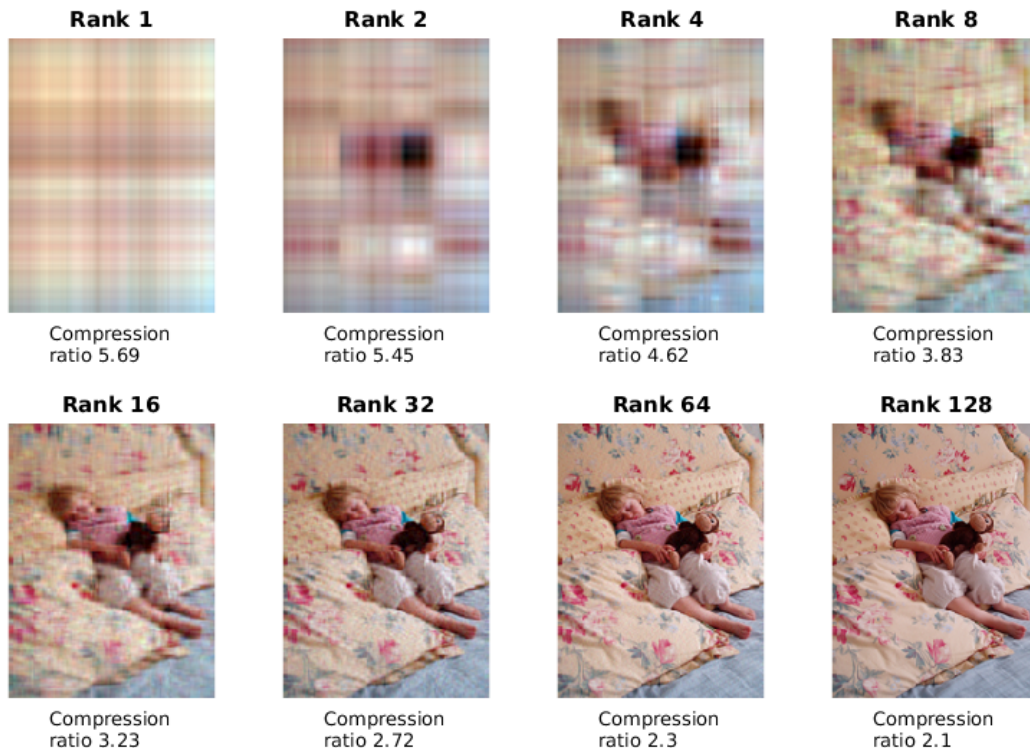
COMPONENTS = 3;
FILE_NAME = '/nena_compressed_rank_';

%Space allocation
compressed_image = uncompressed_image;
compressed_image_divided_by_largest = uncompressed_image;
error_sum = 0;

for i = 1:COMPONENTS
    uncompressed_component = uncompressed_image(:,:,i);
    compressed_component = svd_compression(uncompressed_component, rank);
    compressed_image_divided_by_largest(:,:,i) = ...
        divide_all_by_largest_element(compressed_component);
    compressed_image(:,:,i) = compressed_component;

    error_sum = error_sum + matrix_relative_error( ...
        uncompressed_component, compressed_component);
end
```

Nena color, SVD compression




```

error_median = error_sum / COMPONENTS;

file_name = create_image_path(FOLDER, strcat(FILE_NAME, int2str(rank)));
imwrite(compressed_image_divided_by_largest, file_name);

size_ratio = compression_ratio(uncompressed_size, file_name);

```

Looking at the table below, it appears that SVD gives better results for color images than for grayscale, with higher compression ratios at every rank and similar errors. Compression ratio 2 is now achieved even at rank 2^7 , a huge improvement over grayscale. This image also has good quality even when zoomed in.

Otherwise, the relationship between compression ranks and image quality matches that of the grayscale model pretty closely, a fact verified by the table. My naked-eye tells me that relative errors of 0.005 and below are generally to be preferred.

Finally, to create a video from the compressed images `assignment_2_video` is called. I experimented with many different values for RANKS and found that linear values of 1:3:100 gave the most interesting flow of images. The variable F is assigned a frame every time a compressed color image is created in `assignment_2_color`, and then passed to this function along with RANKS.

```

function [] = assignment_2_video(F, RANKS)

FRAME_RATE = 6;

RANKS_FIRST = int2str(RANKS(1));
RANKS_DIFF = int2str(RANKS(2) - RANKS(1));
RANKS_LAST = int2str(RANKS(end));
RANKS_FOR_TITLE = [RANKS_FIRST, ':', RANKS_DIFF, ':', RANKS_LAST];
VIDEO_PATH = [VIDEO_FOLDER, TITLE_START, RANKS_FOR_TITLE, ...
    '_framerate_', int2str(FRAME_RATE), '.avi'];

myVideo = VideoWriter(VIDEO_PATH);
myVideo.FrameRate = FRAME_RATE;

open(myVideo);
writeVideo(myVideo, F);
close(myVideo);

```

A possible error source throughout the program is the manipulation of images. Rescaling the subplot figures in Matlab sometimes give rise to clear artifacts, although this can be (mostly) eliminated by manual resizing to satisfaction before saving. I saved the subplot images in the lossless .bmp format for importing to this report, and then avoided further rescaling them in the text editor.

For printing the report I exported it as a .pdf and made sure to use a lossless compression for the images. However, the actual process of printing on paper probably adds a certain amount of distortion all images.

From this project I learned huge amounts about Matlab, like how to transform data for your needs, how to manipulate strings in useful ways and how to work with files and folders in a script. I learned how to write better Matlab functions, processing images and how to export data, including into video. I also got better acquainted with how images are represented internally, which was quite eye-opening.

Working with some functions that only accept 8-bit integer matrices, and others that only accept numbers in the range [0,1] was frustrating at first. Same thing with the video functionality where the Matlab help pages are not always easy to read, although the Matlab Central forums are usually more accommodating.

Going through the painstaking process of finding the right function and syntax for each situation was still probably a worthwhile endeavor, and I already feel much more comfortable with using Matlab as a proper scripting language. It has tons of power.

SVD compression rank	Nena grayscale relative error	Nena color relative error
1	0.14	0.14
2	0.099	0.10
4	0.057	0.061
8	0.036	0.038
16	0.024	0.024
32	0.013	0.014
64	0.0073	0.0077
128	0.0036	0.0038
256	0.0016	0.0017
512	0.00065	0.00069
1024	0.00016	0.00017